

Diffusion model for machine learning

Tran Trong Khiem

AI lab training

2024/08/01

1 Introduction

2 NCSN

3 DDPM

4 References

Introduction

Math for machine learning :

- Complete Foundation chapter in Probabilistic Machine Learning[1]
- Probability and Statistics.
- Linear Algebra.
- Optimazation.

Generative AI:

- Gan
- VAE
- Flow-base
- Diffusion model

Generative AI

Existing **generative modeling techniques** can largely be grouped into two categories based on how they **represent probability distributions**.

- ① **likelihood-based models**: which **directly learn the distribution's** probability density (or mass) function via (approximate) **maximum likelihood**. (VAEs, EBMs, ...)
 - **Cons**: require **strong restrictions on the model architecture** to ensure a tractable normalizing constant for **likelihood computation**.
- ② **implicit generative models**: where the **probability distribution** is implicitly represented by a **model of its sampling process**. (Gan,...)
 - **Cons**: **unstable** and can lead to **model collapse**.

Diffusion model introduces **another way** to represent **probability distributions** that **circumvent several of these limitations**.

Diffusion model

The **key idea** is to **model the gradient of the log probability density function**, score function.

- **score-based models** are not required to have a tractable normalizing constant, and can be directly learned by **score matching**. Better than GAN in image generation.

Denote :

- The dataset consists of i.i.d. samples $\{x_i \in \mathbb{R}^D\}_{i=1}^N$ from an **unknown data distribution** $p_{\text{data}}(x)$.
- The **score** of a probability density $p(x)$ is defined as $\nabla_x \log p(x)$.
- The score network $s_\theta : \mathbb{R}^D \rightarrow \mathbb{R}^D$, which will be **trained to approximate** the score of $p_{\text{data}}(x)$.

The **framework of score-based generative modeling**:

- ① score matching
- ② Langevin dynamics.

Framework of score-based generative modeling

Score matching :

- train a **score network** $s_\theta(x)$ to estimate $\nabla_x \log p_{\text{data}}(x)$ without training a model to estimate $p_{\text{data}}(x)$

Langevin dynamics

- produce samples from a probability density $p(x)$ **using only the score function** $\nabla_x \log p_{\text{data}}(x)$.




Figure 1: Score-based generative modeling with score matching + Langevin dynamics.

Score matching for score estimation

Goal: train a **score network** $s_\theta(x)$ to estimate $\nabla_x \log p_{\text{data}}(x)$.
The objective minimizes :

$$\mathbb{E}_{p_{\text{data}}} \left[\|s_\theta(x) - \nabla_x \log p_{\text{data}}(x)\|_2^2 \right]$$

which can be shown equivalent to the following up to a constant :

$$\mathbb{E}_{p_{\text{data}}(x)} \left[\text{tr}(\nabla_x s_\theta(x)) + \frac{1}{2} \|s_\theta(x)\|_2^2 \right]$$

Problem: Score matching is **not scalable to deep networks and high-dimensional data** due to the computation of $\text{tr}(\nabla_x s_\theta(x))$.

Solution: There are two popular methods for large scale score matching.

- 1 Denoising score matching
- 2 Sliced score matching

Score matching for score estimation(.cnt)

Denoising score matching:

- completely circumvents $\text{tr}(\nabla_x s_\theta(x))$.
- perturbs the data point x with a prespecified noise $q_\sigma(\tilde{x} | x)$.
- employs score matching to **estimate the score of the perturbed data**.

$$\mathbb{E}_{q_\sigma(\tilde{x}|x)p_{\text{data}}(x)} \left[\|s_\theta(\tilde{x}) - \nabla_{\tilde{x}} \log q_\sigma(\tilde{x} | x)\|_2^2 \right]$$

- However, $s_\theta^*(x) = \nabla_x \log q_\sigma(x) \approx \nabla_x \log p_{\text{data}}(x)$ is true only when **the noise is small enough** such that $q_\sigma(x) \approx p_{\text{data}}(x)$.

Sliced score matching:

- uses random projections to approximate $\text{tr}(\nabla_x s_\theta(x))$.
- The objective is:

$$\frac{1}{2} \left| \mathbb{E}_{p_v} \mathbb{E}_{p_{\text{data}}} \left[v \nabla_x s_\theta(x) v + \|s_\theta(x)\|_2^2 \right] \right|$$

- p_v is a simple distribution of random vectors.

Sampling with Langevin dynamics

Goal: produce samples from a probability density $p(x)$ using only the score function $\nabla_x \log p(x)$.

- Given a fixed step size $\epsilon > 0$, and an initial value $\tilde{x}_0 \sim \pi(x)$
- π is a prior distribution.
- Langevin method recursively computes the following :

$$\tilde{x}_t = \tilde{x}_{t-1} + \frac{\epsilon}{2} \nabla_x \log p(\tilde{x}_{t-1}) + \epsilon z_t,$$

- $z_t \sim \mathcal{N}(0, I)$
- The distribution of \tilde{x}_T equals $p(x)$ when $\epsilon \rightarrow 0$ and $T \rightarrow \infty$,
- In practice, ϵ is small and T is large.

Challenges of score-based generative modeling

Inaccurate score estimation with score matching:

- In score matching, we minimize :

$$\mathbb{E}_{p_{\text{data}}} \left[\|s_{\theta}(x) - \nabla_x \log p_{\text{data}}(x)\|_2^2 \right] = \int p(x) \left[\|s_{\theta}(x) - \nabla_x \log p_{\text{data}}(x)\|_2^2 \right] dx$$

- Since square error weighted by $p(x)$, they are largely **ignored in low density regions** where $p(x)$ is small.




Figure 2: Estimated scores are only accurate in high density regions

How to bypass the inaccurate score estimation in regions of low data density?

Observation : perturbing data with random Gaussian noise makes the data distribution more **amenable to score-based generative modeling**.

- **large Gaussian noise** has the effect of **filling low density regions** in the original distribution.

Upon intuition is the key idea for **Noise Conditional Score Networks(NCSN)**:

- ① perturbing the data using various levels of noise.
- ② simultaneously estimating scores corresponding to all noise levels by training a single conditional score network.

1 Introduction

2 NCSN

3 DDPM

4 References

Noise Conditional Score Networks

Problem : How to choose an **appropriate noise scale** for the perturbation process?

- Larger noise over-corrupts the data and alters it significantly from the original distribution.
- Smaller noise, on the other hand, causes less corruption of the original data.

Solution: Use **multiple scales of noise** perturbations simultaneously.
Denote:

- $\{\sigma_i\}_{i=1}^L$ be a positive sequence geometric decending sequence.
- $q_\sigma(x) = \int p_{\text{data}}(t) \mathcal{N}(x | t, \sigma^2 I) dt$ the **perturbed data distribution**.
- $s_\theta(x, \sigma)$ is a Noise Conditional Score Network (NCSN).
- train model to jointly estimate the scores of **all perturbed data distributions** :

$$\forall \sigma \in \{\sigma_i\}_{i=1}^L : s_\theta(x, \sigma) \approx \nabla_x \log q_\sigma(x)$$

Learning NCSNs via score matching

Adapt **denoising score matching** for learning NCSNs.

- choose the noise distribution to be $q_\sigma(\tilde{x} | x) = \mathcal{N}(\tilde{x} | x, \sigma^2 I)$
- therefore $\nabla_{\tilde{x}} \log q_\sigma(\tilde{x} | x) = -\frac{\tilde{x} - x}{\sigma^2}$
- For a given σ , the denoising score matching objective is :

$$\mathcal{L}(\theta; \sigma) = \frac{1}{2} \mathbb{E}_{p_{\text{data}}(x)} \mathbb{E}_{\tilde{x} \sim \mathcal{N}(x, \sigma^2 I)} \left[\left\| s_\theta(\tilde{x}, \sigma) + \frac{\tilde{x} - x}{\sigma^2} \right\|_2^2 \right].$$

- We combine for all $\sigma \in \{\sigma_i\}_{i=1}^L$ to get one unified objective :

$$\mathcal{L}(\theta; \{\sigma_i\}_{i=1}^L) = \frac{1}{L} \sum_{i=1}^L \lambda(\sigma_i) \mathcal{L}(\theta; \sigma_i)$$

NCSN inference via annealed Langevin dynamics

- propose a sampling approach— **annealed Langevin dynamics**

Algorithm 1 Annealed Langevin dynamics.

Require: $\{\sigma_i\}_{i=1}^L, \epsilon, T$.

1: Initialize $\tilde{\mathbf{x}}_0$

2: **for** $i \leftarrow 1$ to L **do**

3: $\alpha_i \leftarrow \epsilon \cdot \sigma_i^2 / \sigma_L^2$ $\triangleright \alpha_i$ is the step size.

4: **for** $t \leftarrow 1$ to T **do**

5: Draw $\mathbf{z}_t \sim \mathcal{N}(0, I)$

6: $\tilde{\mathbf{x}}_t \leftarrow \tilde{\mathbf{x}}_{t-1} + \frac{\alpha_i}{2} \mathbf{s}_\theta(\tilde{\mathbf{x}}_{t-1}, \sigma_i) + \sqrt{\alpha_i} \mathbf{z}_t$

7: **end for**

8: $\tilde{\mathbf{x}}_0 \leftarrow \tilde{\mathbf{x}}_T$

9: **end for**

return $\tilde{\mathbf{x}}_T$

Figure 3: Annealed Langevin dynamics.

1 Introduction

2 NCSN

3 DDPM

4 References

Denoising Diffusion Probabilistic Models




Figure 4: Diffusion model

Forward diffusion process:

- add small amount of Gaussian noise to the sample in T
- producing a sequence of noisy samples $x_1, x_2 \dots x_T$
- converts any complex data distribution into a simple, tractable, distribution.

Reverse diffusion process:

- Learn a reversal of forward diffusion process.

Foward process

Gradually adds **Gaussian noise** to the data according to a variance schedule β_1, \dots, β_T :


$$q(x_{1:T} | x_0) := \prod_{t=1}^T q(x_t | x_{t-1}), \quad q(x_t | x_{t-1}) := \mathcal{N}(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$$

Nice property: We can sample x_t at timestep t as :

$$x_t = \sqrt{\alpha_t}x_{t-1} + \sqrt{1 - \alpha_t}\epsilon_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon$$

- $\epsilon_t \sim \mathcal{N}(0, I)$
- $\bar{\alpha}_t = \prod_{s=1}^t \alpha_s$ and $\alpha_t = 1 - \beta_t$
- Thus we have : $q(x_t|x_0) = \mathcal{N}(x_t, \sqrt{\bar{\alpha}_t}x_0, (1 - \bar{\alpha}_t)I)$

Reverse diffusion process



Goal: Learn to reverse the **forward process** and sample from $q(x_{t-1}|x_t)$.

- Use $p_\theta(x_{t-1}|x_t)$ to approximate $q(x_{t-1}|x_t)$.
- The reverse conditional probability is **tractable when conditioned on x_0** :

$$q(x_{t-1}|x_t, x_0) = \mathcal{N}(x_{t-1}, \tilde{\mu}(x_t, x_0), \tilde{\beta}_t I)$$

- $\tilde{\beta}_t = \frac{1-\alpha_{t-1}}{1-\alpha_t} \beta_t$ and $\tilde{\mu}(x_t, x_0) = \sqrt{\alpha_t} \left(\frac{1-\alpha_{t-1}}{1-\alpha_t} \right) x_t + \left(\frac{\sqrt{\alpha_{t-1}} \beta_t}{1-\alpha_t} \right) x_0$

Reverse diffusion process(.cnt)

Training is performed by **optimizing the usual variational bound** on negative log likelihood:

$$\begin{aligned}\mathbb{E}_q[-\log p_\theta(x_0)] &\leq \mathbb{E}_q \left[-\log \frac{p_\theta(x_{0:T})}{q(x_{1:T} | x_0)} \right] \\ &= \mathbb{E}_q \left[-\log p(x_T) - \sum_{t=1}^T \log \frac{p_\theta(x_{t-1} | x_t)}{q(x_t | x_{t-1})} \right] \quad =: L.\end{aligned}$$

Loss function can rewrite as :

$$\mathbb{E}_q \left[D_{KL}(q(x_T|x_0)||p(x_T)) + \sum_{t>1} D_{KL}(q(x_{t-1}|x_t, x_0)||p_\theta(x_{t-1}|x_t)) - \log p_\theta(x_0|x_1) \right] \quad (1)$$

Label each component in the variational lower bound loss separately:

- $L_{VLB} = \sum_{t=0}^T L_t$

Reverse diffusion process(.cnt)

The loss term L_t is parameterized and simplified to minimize :

$$L_t^{simple} = \mathbb{E}_{t \sim [1, T], \mathbf{x}_0, \epsilon_t} \left[\|\epsilon_t - \epsilon_\theta(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon_t t)\|^2 \right]$$

Algorithm 1 Training

```

1: repeat
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$ 
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$ 
4:    $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
5:   Take gradient descent step on
        $\nabla_\theta \|\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t)\|^2$ 
6: until converged

```

Algorithm 2 Sampling

```

1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$ 
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \alpha_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$ 
5: end for
6: return  $\mathbf{x}_0$ 

```

Figure 5: Traing process.

1 Introduction

2 NCSN

3 DDPM

4 References

References

- ① Weng, Lilian. (Jul 2021). What are diffusion models? Lil'Log.
- ② Jonathan Ho, Ajay Jain, Pieter Abbeel, Denoising Diffusion Probabilistic Models
- ③ Yang Song, Stefano Ermon, Generative Modeling by Estimating Gradients of the Data Distribution
- ④ Generative Modeling by Estimating Gradients of the Data Distribution